

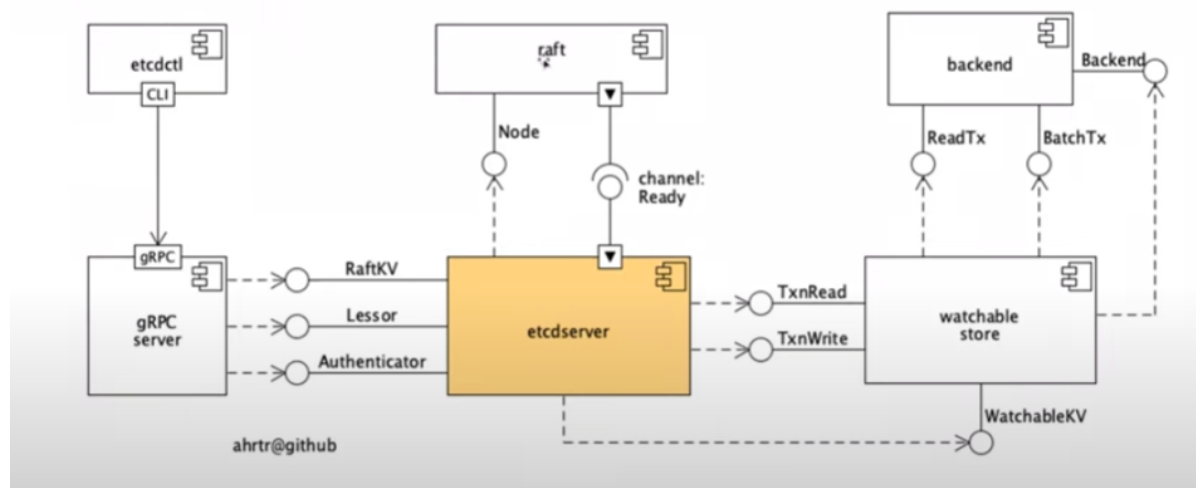
Etcd Deep Dive

Deep view of etcd from [this video](#).

1 General architecture

- etcd is a distributed reliable key-value store for the most critical data of a distributed system.
 - Key-value
 - It's a No-SQL database.
 - Strong consistency.
 - etcd is a CP system.
 - It can't serve write requests if the quorum isn't satisfied.
 - It support both linearizable (default) and serializable read requests.
 - `etcdctl get k1 --consistency="s"`
 - `curl http://127.0.0.1:2379/health?serializable=true`

Architecture



- Read: gRPC server -> RaftKV -> etcdserver --(propose) --> Raft --(commit)-->etcdserver -> store

2 Storage: 3 kinds

2-1 WAL

video: 19:13

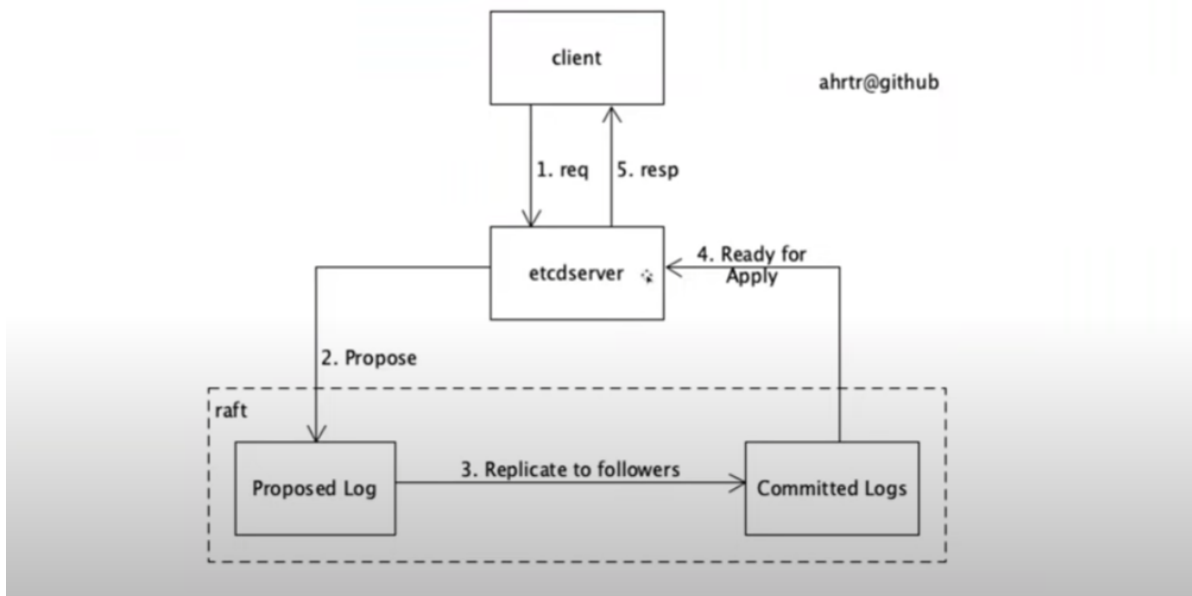
2-2

2-3

3 Raft replication

- Communication between raft and etcdserver
 - etcdserver talks to raft via Node interface.
 - Raft talks to etcdserver via channel.

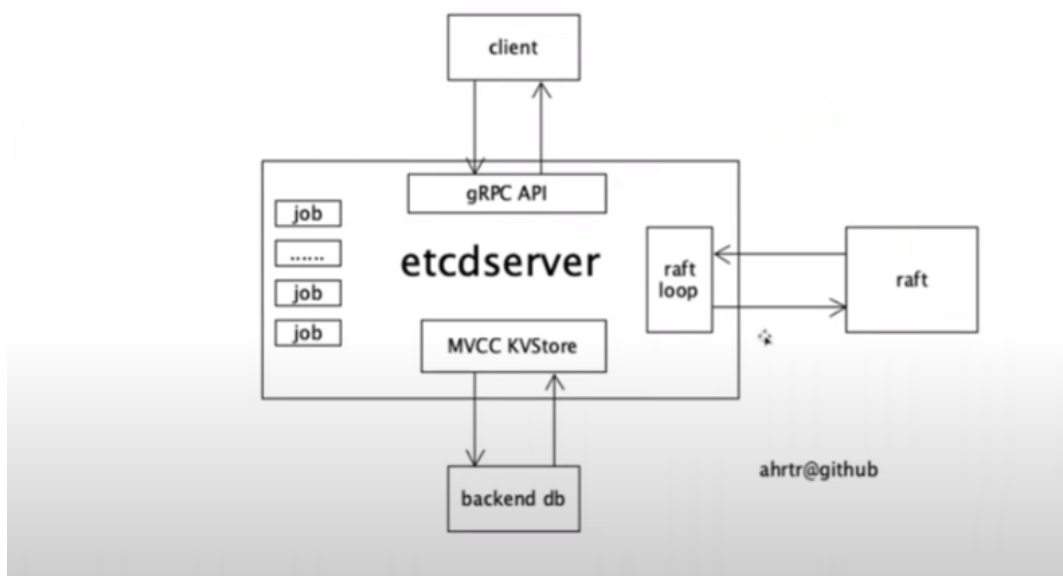
Replication: Raft



- Write: all go through the leader
- Read:
 - serializable: go through leader or follower
 - linerizable: only go through leader

4 Core server

etcdserver



- **Core:** gRPC && Raft

5 A PUT workflow

Put e2e workflow

